


**DEPARTMENT OF COMPUTER SCIENCE AND  
ENGINEERING**

**GOVERNMENT ENGINEERING COLLEGE, THRISSUR**




**Practical Record**  
**221LCS100 - COMPUTING LAB 1**

**AROMAL M B**  
**TCR24CSC04**

# **DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING**

**GOVERNMENT ENGINEERING COLLEGE, THRISSUR**



## **CERTIFICATE**

This is to certify that this is a bonafide report of the work done in **221LCS100 - COMPUTING LAB 1** by **Mr. AROMAL M B (TCR24CSCE04)** under our guidance towards partial fulfillment of the requirement for the award of the Degree of Master of Technology in Computer Science and Engineering of APJ Abdul Kalam Technological University during the year 2024 - 2026.

**Prof. Dr. Giles M P,**  
Professor  
Dept. of CSE,  
Govt Engineering College, Thrissur

**Prof. Dr. Ajay James,**  
Associate Professor and HOD,  
Dept. of CSE,  
Govt Engineering College, Thrissur

# Contents

<b>Decision Tree - ID3 Algorithm</b>	<b>1</b>
<b>Naive Bayes Classifier on Iris Dataset</b>	<b>4</b>
<b>Naive Bayes Classifier on Iris Dataset (with Evaluation Metrics)</b>	<b>6</b>
<b>Bayesian Network for Heart Disease Diagnosis</b>	<b>10</b>
<b>Clustering Using EM Algorithm and K-Means</b>	<b>14</b>
<b>Classification Using k-Nearest Neighbour Algorithm</b>	<b>19</b>
<b>Locally Weighted Regression (LWR)</b>	<b>23</b>
<b>5-Fold Cross Validation</b>	<b>26</b>
<b>Support Vector Machine for CIFAR-10 Classification</b>	<b>29</b>
<b>Artificial Neural Network with Backpropagation</b>	<b>33</b>

# Experiment 1

## Decision Tree - ID3 Algorithm

### Aim

The objective of this experiment is to implement the ID3 algorithm for constructing decision trees using information gain.

### Algorithm

1. **Load Libraries:** Import necessary libraries: `pandas`, `numpy`, `matplotlib`, `LabelEncoder`, `DecisionTreeClassifier`, and `train_test_split`.
2. **Load Dataset:** Read the dataset `tennis.csv` into a `DataFrame`.
3. **Encode Data:** Convert categorical data to numeric using `LabelEncoder`.
4. **Split Data:** Separate features (X) and target (Y), then split into training and testing sets using `train_test_split`.
5. **Train Model:** Initialize and train a `DecisionTreeClassifier` with `entropy` criterion.
6. **Make Predictions:** Predict outcomes on the test data.
7. **Evaluate Accuracy:** Print the model's accuracy on the test set.
8. **Single Prediction:** Make a prediction for a sample input.
9. **Visualize Tree:** Plot the decision tree.

# Implementation

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from sklearn.preprocessing import LabelEncoder
5 from sklearn.tree import DecisionTreeClassifier, plot_tree
6 from sklearn.model_selection import train_test_split
7
8 # Load the dataset
9 df = pd.read_csv('/content/tennis.csv')
10
11 # Initialize LabelEncoder
12 encoder = LabelEncoder()
13
14 # Encode all categorical columns
15 for col in df.columns:
16     df[col] = encoder.fit_transform(df[col])
17
18 # Split data into features and target
19 X = df.drop(columns=['play']) # Features
20 Y = df['play'] # Target
21
22 # Split data into training and testing sets
23 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
24
25 # Train the Decision Tree model
26 model = DecisionTreeClassifier(criterion='entropy', random_state=42)
27 model.fit(X_train, Y_train)
28
29 # Make predictions and evaluate accuracy
30 Y_pred = model.predict(X_test)
31 accuracy = model.score(X_test, Y_test)
32 print(f"\nModel Accuracy: {accuracy * 100:.2f}%")
33 Model Accuracy: 100.00%
34
35 # Example prediction
36 sample_input = [[2, 1, 0, 1]] # Example input: [outlook, temp, humidity, windy]
37 prediction = model.predict(sample_input)
38 print(f"\nPrediction for input {sample_input}: {prediction} (1 = Play, 0 = Don't Play)")
39 Prediction for input [[2, 1, 0, 1]]: [0] (1 = Play, 0 = Don't Play)
40
41 # Visualize the Decision Tree
42 plt.figure(figsize=(9,9))
43 plot_tree(model, feature_names=X.columns, class_names=encoder.classes_, filled=True)
44 plt.show()

```

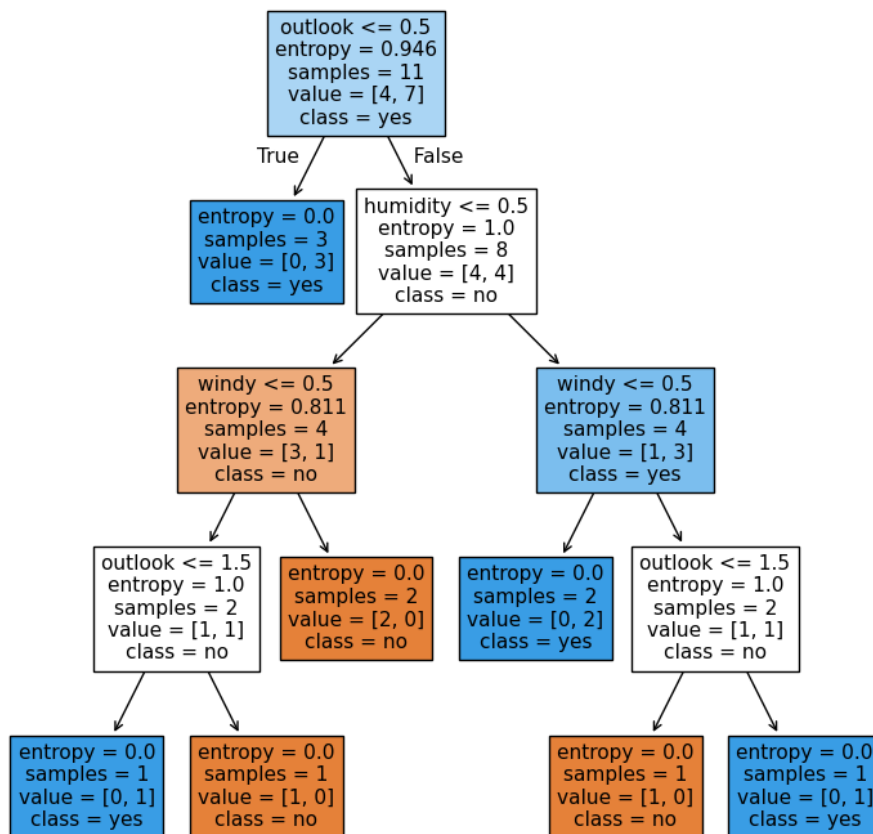


Figure 1: Decision Tree

## Result

Successfully completed the program by implementing the Decision Tree (ID3) algorithm for classifying the tennis dataset.

# Experiment 2

## Naive Bayes Classifier on Iris Dataset

### Aim

Write a program to implement the naïve bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering few test data sets.

### Algorithm

1. **Load Libraries:** Import necessary libraries: `pandas`, `numpy`, and Naive Bayes from `sklearn`.
2. **Load Dataset:** Load the Iris dataset from `Iris.csv`.
3. **Preprocessing:** Drop the 'Id' column and ensure that all features are numerical.
4. **Split Data:** Split the dataset into features (X) and target (y), and then split the data into training and testing sets.
5. **Train Model:** Initialize and train a `GaussianNB` model using the training data.
6. **Make Predictions:** Use the trained model to predict the target for the test data.
7. **Evaluate Accuracy:** Calculate and print the accuracy of the model on the test set.
8. **Single Prediction:** Make predictions for a few sample inputs.

# Implementation

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from sklearn.model_selection import train_test_split
6 from sklearn.naive_bayes import GaussianNB
7 from sklearn.metrics import accuracy_score
8
9 # Load the Iris dataset
10 iris = pd.read_csv('/content/Iris.csv')
11
12 # Drop the 'Id' column
13 iris.drop(columns="Id", inplace=True)
14
15 # Define features and target
16 X = iris.drop(columns='Species')
17 y = iris['Species']
18
19 # Split the dataset into training and testing sets
20 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=0)
21 print("\nShapes of datasets:")
22 print(f"X_train: {X_train.shape}, X_test: {X_test.shape}")
23 print(f"y_train: {y_train.shape}, y_test: {y_test.shape}")
24 Shapes of datasets:
25 X_train: (105, 4), X_test: (45, 4)
26 y_train: (105,), y_test: (45,)
27
28 # Train the Gaussian Naive Bayes model
29 model = GaussianNB()
30 model.fit(X_train, y_train)
31
32 # Make predictions
33 y_pred = model.predict(X_test)
34
35 # Evaluate accuracy
36 accuracy = accuracy_score(y_test, y_pred)
37 print(f"\nAccuracy of the Naive Bayes model: {accuracy * 100:.2f}%")
38 Accuracy of the Naive Bayes model: 100.00%
39
40 # Example predictions
41 X_test_example = [
42     [5.1, 3.5, 1.4, 0.2],
43     [4.9, 3.0, 1.4, 0.2],]
44 y_pred_example = model.predict(X_test_example)
45 print("Predictions:", y_pred_example)
46 Predictions: ['Iris-setosa' 'Iris-setosa']

```

## Result

Successfully implemented the Naive Bayes classifier on the Iris dataset.



# Experiment 3

## Naive Bayes Classifier on Iris Dataset (with Evaluation Metrics)

### Aim

The objective of this experiment is to apply the Naive Bayes classifier on the Iris dataset, evaluate its performance using accuracy, precision, recall, and F1-score, and visualize class-wise metrics.

### Algorithm

1. **Load Libraries:** Import the necessary libraries: `pandas`, `sklearn`, `matplotlib`, `seaborn`.
2. **Load Dataset:** Read the Iris dataset from `Iris.csv`.
3. **Preprocess Data:** Separate features (X) and target (y). Split the data into training and testing sets.
4. **Standardize Data:** Standardize the features using `StandardScaler`.
5. **Train Model:** Train a `GaussianNB` model using the training data.
6. **Make Predictions:** Predict the target values for the test set.
7. **Evaluate Model:** Calculate and print accuracy, precision, recall, and F1-score.
8. **Prediction for New Sample:** Make predictions for new sample data points.
9. **Visualize Metrics:** Visualize the precision, recall, and F1-score for each class.

# Implementation

```

1 import pandas as pd
2 from sklearn.model_selection import train_test_split
3 from sklearn.preprocessing import StandardScaler
4 from sklearn.naive_bayes import GaussianNB
5 from sklearn.metrics import classification_report, accuracy_score, \
6     precision_score, recall_score, f1_score
7 import matplotlib.pyplot as plt
8 import seaborn as sns
9
10 # Load the Iris dataset
11 file_path = '/content/Iris.csv'
12 data = pd.read_csv(file_path)
13
14 # Separate features and target
15 X = data[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']] # Features
16 y = data['Species'] # Target variable
17
18 # Split the dataset into training and testing sets
19 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
20
21 # Standardize the features
22 scaler = StandardScaler()
23 X_train = scaler.fit_transform(X_train)
24 X_test = scaler.transform(X_test)
25
26 # Train the Gaussian Naive Bayes model
27 nb_model = GaussianNB()
28 nb_model.fit(X_train, y_train)
29
30 # Make predictions
31 y_pred = nb_model.predict(X_test)
32
33 # Calculate evaluation metrics
34 accuracy = accuracy_score(y_test, y_pred)
35 precision = precision_score(y_test, y_pred, average='weighted')
36 recall = recall_score(y_test, y_pred, average='weighted')
37 Accuracy: 0.9777777777777777
38 Precision: 0.9793650793650793
39 Recall: 0.9777777777777777
40
41 # Print metrics
42 print("Accuracy:", accuracy)
43 print("Precision:", precision)
44 print("Recall:", recall)
45
46 # Example prediction for a new sample
47 new_sample = [[5.0, 3.2, 4.0, 1.5]]
48 new_prediction = nb_model.predict(new_sample)
49 print("Prediction for new sample:", new_prediction[0])
50 Prediction for new sample: Iris-virginica
51
52 # Print classification report
53 print("\nClassification Report:")
54 print(classification_report(y_test, y_pred))
55
56 Classification Report:
57           precision    recall  f1-score   support

```

```

58
59      Iris-setosa      1.00      1.00      1.00      19
60 Iris-versicolor      1.00      0.92      0.96      13
61 Iris-virginica       0.93      1.00      0.96      13
62
63      accuracy                0.98      45
64      macro avg      0.98      0.97      0.97      45
65      weighted avg    0.98      0.98      0.98      45
66
67
68 # Calculate precision, recall, and F1-score by class
69 precision = precision_score(y_test, y_pred, average=None, labels=nb_model.classes_)
70 recall = recall_score(y_test, y_pred, average=None, labels=nb_model.classes_)
71 f1 = f1_score(y_test, y_pred, average=None, labels=nb_model.classes_)
72
73 # Dataframe for class-wise metrics
74 metrics = pd.DataFrame({
75     'Class': nb_model.classes_,
76     'Precision': precision,
77     'Recall': recall,
78     'F1-Score': f1
79 }).melt(id_vars='Class', var_name='Metric', value_name='Value')
80
81 # Plot the metrics
82 plt.figure(figsize=(10, 6))
83 sns.barplot(data=metrics, x='Class', y='Value', hue='Metric')
84 plt.title('Class-wise Precision, Recall, and F1-Score')
85 plt.ylim(0, 1.1)
86 plt.ylabel('Score')
87 plt.xlabel('Class')
88 plt.legend(loc='upper right')
89 plt.show()

```




Figure 2: Class-wise Precision, Recall, and F1-Score

## Result

The Naive Bayes classifier successfully classified the Iris dataset. The model's metrics, including accuracy, precision, recall, and F1-score, were evaluated, and class-wise metrics were visualized.

# **Experiment 4**

## **Bayesian Network for Heart Disease**

### **Diagnosis**

#### **Aim**

To construct a Bayesian Network using medical data and demonstrate the diagnosis of heart patients using the standard Heart Disease dataset.

#### **Algorithm**

1. Import the necessary libraries
2. Load the Heart Disease dataset and preprocess the data by handling missing values and encoding categorical features.
3. Define the structure of the Bayesian Network using domain knowledge or a structure-learning algorithm.
4. Train the Bayesian Network using Maximum Likelihood Estimation (MLE) or other suitable methods.
5. Perform inference on the network:
  - Query probabilities of the target variable ('target') based on given evidence.
  - Predict the presence or absence of heart disease for individual patients.
6. Evaluate the model's performance by comparing predictions with true labels and calculating accuracy.
7. Plot the accuracy of predictions to analyze model performance.

# Implementation

```

1 !pip install pgmpy
2 import pandas as pd
3 from pgmpy.models import BayesianNetwork
4 from pgmpy.estimators import ParameterEstimator,
5 from pgmpy.estimators import MaximumLikelihoodEstimator
6 from pgmpy.inference import VariableElimination
7
8 # Load the dataset
9 file_path = '/content/heart.csv'
10 data = pd.read_csv(file_path)
11
12 # Display dataset information
13 print("Dataset Head:\n", data.head())
14 print("\nDataset Info:\n", data.info())
15
16 # Preprocess the dataset
17 # Example: Assuming columns like 'sex', 'cp', and 'thal' are categorical
18 data[categorical_cols] = data[categorical_cols].astype('category')
19
20 # Define the Bayesian Network structure
21 # Example structure based on domain knowledge
22 model = BayesianNetwork([
23     ('age', 'trestbps'),
24     ('age', 'chol'),
25     ('sex', 'thal'),
26     ('chol', 'target'),
27     ('thal', 'target'),
28     ('trestbps', 'target'),])
29
30 # Train the model using Maximum Likelihood Estimation
31 model.fit(data, estimator=MaximumLikelihoodEstimator)
32
33 # Perform inference
34 inference = VariableElimination(model)
35
36 # Query: Probability of having heart disease given specific conditions
37 query_result = inference.query(
38     variables=['target'],
39     evidence={'sex': 1, 'age': 58, 'thal': 2})
40 print("\nQuery Result (Probability of Heart Disease):")
41 print(query_result)
42
43 # Predict for the entire dataset (optional)
44 predicted_target = []
45 for _, row in data.iterrows():
46     evidence = row[['sex', 'age', 'thal']].to_dict()
47     result = inference.map_query(variables=['target'], evidence=evidence)
48     predicted_target.append(result['target'])
49
50 # Add predictions to the dataset
51 data['predicted_target'] = predicted_target
52 print("\nSample Predictions:")
53 print(data[['age', 'sex', 'thal', 'target', 'predicted_target']].head())
54
55 import matplotlib.pyplot as plt
56 from sklearn.metrics import accuracy_score
57

```

```

58 # Assuming `data['target']` is the true target and
59 # `data['predicted_target']` is the predicted target
60 true_labels = data['target']
61 predicted_labels = data['predicted_target']
62
63 # Calculate overall accuracy
64 accuracy = accuracy_score(true_labels, predicted_labels)
65 print(f"Overall Accuracy: {accuracy:.2f}")
66
67 # Plotting accuracy
68 # Simulate accuracy over sample sizes for demonstration
69 sample_sizes = range(10, len(data), 10)
70 accuracies = [
71     accuracy_score(true_labels[:size], predicted_labels[:size])
72     for size in sample_sizes]
73
74 # Plot accuracy vs. sample size
75 plt.figure(figsize=(10, 6))
76 plt.plot(sample_sizes, accuracies, marker='o', linestyle='-', \
77         color='blue', label='Accuracy')
78 plt.axhline(y=accuracy, color='red', linestyle='--', \
79         label=f'Overall Accuracy: {accuracy:.2f}')
80 plt.title("Accuracy of Bayesian Network Predictions")
81 plt.xlabel("Number of Samples")
82 plt.ylabel("Accuracy")
83 plt.legend()
84 plt.grid()
85 plt.show()
86
87 #Output
88 Query Result (Probability of Heart Disease):
89 +-----+-----+
90 | target | phi(target) |
91 +-----+-----+
92 | target(0) | 0.4821 |
93 +-----+-----+
94 | target(1) | 0.5179 |
95 +-----+-----+
96
97 Sample Inputs
98   ca  thal  target
99 0   2    3      0
100 1   0    3      0
101 2   0    3      0
102 3   1    3      0
103 4   3    2      0
104
105 Prediction on Sample Inputs
106 Sample Predictions:
107   age  sex  thal  target  predicted_target
108 0   52   1    3      0                0
109 1   53   1    3      0                0
110 2   70   1    3      0                0
111 3   61   1    3      0                0
112 4   62   0    2      0                1
113 Overall Accuracy: 0.79

```




Figure 3: Accuracy Curve

## Result

The Bayesian Network was successfully constructed and trained using the Heart Disease dataset. The network accurately predicted the presence or absence of heart disease.



# Experiment 5

## Clustering Using EM Algorithm and K-Means

### Aim

Apply EM algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using k-Means algorithm. Compare the results of these two algorithms and comment on the quality of clustering.

### Algorithm

1. **Import Libraries:** Import necessary libraries
2. **Load Data:** Load the Iris dataset and extract the features.
3. **Preprocessing:** Standardize the dataset to scale the features.
4. **Gaussian Mixture Model (EM Algorithm):**
  - Apply the Gaussian Mixture Model (GMM) with the number of clusters set to 3.
  - Fit the model to the dataset and predict the cluster labels.
5. **K-Means Clustering:**
  - Apply K-Means clustering with 3 clusters.
  - Fit the model and predict the cluster labels.
6. **Evaluate and Compare:** Compare the clustering results using silhouette scores and visualizations.
7. **Visualize Clusters:** Visualize the clusters obtained from both algorithms using scatter plots.

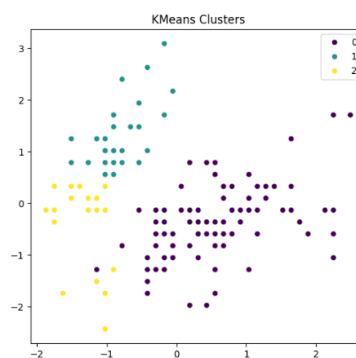
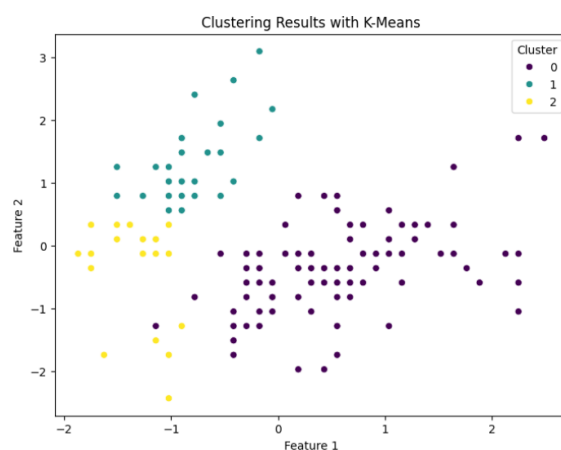
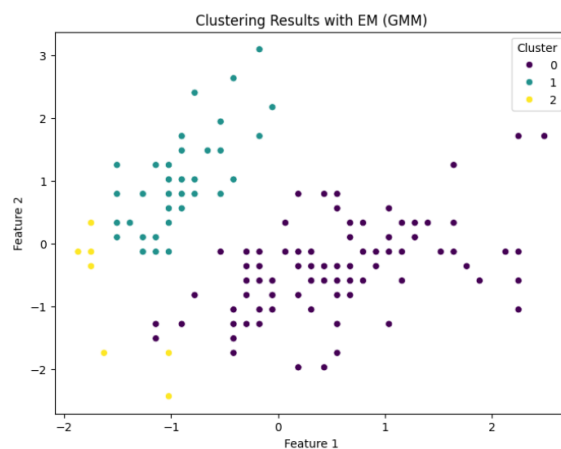
# Implementation

```

1 # Importing necessary libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.mixture import GaussianMixture
7 from sklearn.cluster import KMeans
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.metrics import silhouette_score, accuracy_score
10
11 # Load the Iris dataset from sklearn
12 from sklearn.datasets import load_iris
13
14 # Load the data into a pandas DataFrame
15 iris = load_iris()
16 data = pd.DataFrame(iris.data, columns=iris.feature_names)
17
18 # Display the first few rows of the dataset
19 data.head()
20
21 # Standardize the data to ensure better performance of clustering algorithms
22 scaler = StandardScaler()
23 X_scaled = scaler.fit_transform(X)
24 X_scaled[:5]
25 array([[ -0.90068117,  1.01900435, -1.34022653, -1.3154443 ],
26        [ -1.14301691, -0.13197948, -1.34022653, -1.3154443 ],
27        [ -1.38535265,  0.32841405, -1.39706395, -1.3154443 ],
28        [ -1.50652052,  0.09821729, -1.2833891 , -1.3154443 ],
29        [ -1.02184904,  1.24920112, -1.34022653, -1.3154443 ]])
30
31 # Apply the EM algorithm using GaussianMixture model (GMM)
32 gmm = GaussianMixture(n_components=3, random_state=42)
33 gmm.fit(X_scaled)
34 gmm_labels = gmm.predict(X_scaled)
35 data['GMM_Cluster'] = gmm_labels
36
37 # Visualize the GMM Clusters
38 plt.figure(figsize=(8, 6))
39 sns.scatterplot(x=X_scaled[:, 0], y=X_scaled[:, 1], hue=gmm_labels, palette='viridis')
40 plt.title('Clustering Results with EM (GMM)')
41 plt.xlabel('Feature 1')
42 plt.ylabel('Feature 2')
43 plt.legend(title='Cluster')
44 plt.show()
45
46 # Apply k-Means clustering
47 kmeans = KMeans(n_clusters=3, random_state=42)
48 kmeans.fit(X_scaled)
49 kmeans_labels = kmeans.predict(X_scaled)
50 data['KMeans_Cluster'] = kmeans_labels
51
52 # Visualize the KMeans Clusters
53 plt.figure(figsize=(8, 6))
54 sns.scatterplot(x=X_scaled[:, 0], y=X_scaled[:, 1], hue=kmeans_labels, palette='viridis')
55 plt.title('Clustering Results with K-Means')
56 plt.xlabel('Feature 1')
57 plt.ylabel('Feature 2')

```

```
58 plt.legend(title='Cluster')
59 plt.show()
60
61 # Evaluating the clustering performance using silhouette score
62 gmm_silhouette = silhouette_score(X_scaled, gmm_labels)
63 kmeans_silhouette = silhouette_score(X_scaled, kmeans_labels)
64 Number of clusters in GMM: 3
65 Number of clusters in KMeans: 3
66
67 print(f'Silhouette Score for GMM: {gmm_silhouette:.4f}')
68 print(f'Silhouette Score for KMeans: {kmeans_silhouette:.4f}')
69 Silhouette Score for GMM: 0.4751
70 Silhouette Score for KMeans: 0.4799
```



## **Result**

Successfully implemented the Gaussian Mixture Model (EM algorithm) and K-Means clustering algorithms on the Iris dataset. The silhouette scores were similar for both.

# Experiment 6

## Classification Using k-Nearest Neighbour Algorithm

### Aim

Implement the k-Nearest Neighbour (k-NN) algorithm to classify the Iris dataset and print both correct and wrong predictions.

### Algorithm

1. **Import Libraries:** Import necessary libraries
2. **Load Data:** Load the Iris dataset from 'sklearn'.
3. **Split Data:** Divide the dataset into training and testing sets with an 80-20 split.
4. **Preprocessing:** Standardize the dataset using 'StandardScaler' for better performance.
5. **Train k-NN Model:**
  - Initialize the k-NN classifier with the desired value of 'k'.
  - Train the model using the scaled training data.
6. **Make Predictions:**
  - Predict the labels for the test dataset.
  - Compare predicted labels with actual labels to identify correct and wrong predictions.
7. **Evaluate:** Calculate classification metrics, accuracy, and display predictions.

# Implementation

```

1 # Importing necessary libraries
2 import numpy as np
3 from sklearn.datasets import load_iris
4 from sklearn.model_selection import train_test_split
5 from sklearn.preprocessing import StandardScaler
6 from sklearn.neighbors import KNeighborsClassifier
7 from sklearn.metrics import classification_report, accuracy_score
8
9 # Load the Iris dataset
10 iris = load_iris()
11 X, y = iris.data, iris.target
12
13 # Split the data into training and testing sets
14 X_train, X_test, y_train, y_test = train_test_split(
15     X, y, test_size=0.2, random_state=42, stratify=y)
16
17 # Standardize the feature values
18 scaler = StandardScaler()
19 X_train = scaler.fit_transform(X_train)
20 X_test = scaler.transform(X_test)
21
22 # Train the k-NN classifier
23 k = 3
24 knn = KNeighborsClassifier(n_neighbors=k)
25 knn.fit(X_train, y_train)
26
27 # Make predictions
28 y_pred = knn.predict(X_test)
29
30 # Evaluate the predictions
31 correct_predictions = [
32     (i, true, pred) for i, (true, pred) in enumerate(zip(y_test, y_pred)) if true == pred]
33 wrong_predictions = [
34     (i, true, pred) for i, (true, pred) in enumerate(zip(y_test, y_pred)) if true != pred]
35
36 # Print correct predictions
37 print("\nCorrect Predictions:")
38 for index, true, pred in correct_predictions:
39     print(f"Index: {index}, True Label: {true}, Predicted Label: {pred}")
40
41 Correct Predictions:
42 Index: 0, True Label: 1, Predicted Label: 1
43 Index: 2, True Label: 0, Predicted Label: 0
44 Index: 3, True Label: 2, Predicted Label: 2
45 Index: 4, True Label: 1, Predicted Label: 1
46 Index: 5, True Label: 1, Predicted Label: 1
47 Index: 6, True Label: 1, Predicted Label: 1
48 Index: 7, True Label: 0, Predicted Label: 0
49 Index: 8, True Label: 0, Predicted Label: 0
50 Index: 9, True Label: 0, Predicted Label: 0
51 Index: 10, True Label: 1, Predicted Label: 1
52 Index: 11, True Label: 0, Predicted Label: 0
53 Index: 12, True Label: 0, Predicted Label: 0
54 Index: 13, True Label: 2, Predicted Label: 2
55 Index: 14, True Label: 2, Predicted Label: 2
56 Index: 15, True Label: 2, Predicted Label: 2
57 Index: 16, True Label: 2, Predicted Label: 2
58 Index: 17, True Label: 2, Predicted Label: 2

```

```

58 Index: 18, True Label: 1, Predicted Label: 1
59 Index: 19, True Label: 0, Predicted Label: 0
60 Index: 20, True Label: 1, Predicted Label: 1
61 Index: 21, True Label: 2, Predicted Label: 2
62 Index: 22, True Label: 1, Predicted Label: 1
63 Index: 23, True Label: 2, Predicted Label: 2
64 ...
65 Index: 26, True Label: 0, Predicted Label: 0
66 Index: 27, True Label: 1, Predicted Label: 1
67 Index: 28, True Label: 2, Predicted Label: 2
68 Index: 29, True Label: 0, Predicted Label: 0
69
70 # Print wrong predictions
71 print("\nWrong Predictions:")
72 for index, true, pred in wrong_predictions:
73     print(f"Index: {index}, True Label: {true}, Predicted Label: {pred}")
74 Wrong Predictions:
75 Index: 1, True Label: 2, Predicted Label: 1
76
77 # Print classification report and accuracy
78 print("\nClassification Report:")
79 print(classification_report(y_test, y_pred, target_names=iris.target_names))
80 print(f"Accuracy: {accuracy_score(y_test, y_pred):.2f}")
81 Classification Report:
82         precision    recall  f1-score   support
83
84     setosa         1.00      1.00      1.00         10
85  versicolor      0.91      1.00      0.95         10
86   virginica      1.00      0.90      0.95         10
87
88     accuracy                   0.97         30
89    macro avg       0.97      0.97      0.97         30
90 weighted avg       0.97      0.97      0.97         30
91
92 Accuracy: 0.97
93
94 # Test the algorithm with new examples
95 def test_knn(example):
96     example_scaled = scaler.transform([example])
97     prediction = knn.predict(example_scaled)
98     predicted_class = iris.target_names[prediction[0]]
99     return predicted_class
100 example_1 = [5.1, 3.5, 1.4, 0.2]
101 example_2 = [6.7, 3.0, 5.2, 2.3]
102 print("\nTesting new examples:")
103 print(f"Example 1: {example_1} -> Predicted Class: {test_knn(example_1)}")
104 print(f"Example 2: {example_2} -> Predicted Class: {test_knn(example_2)}")
105
106 Testing new examples:
107 Example 1: [5.1, 3.5, 1.4, 0.2] -> Predicted Class: setosa
108 Example 2: [6.7, 3.0, 5.2, 2.3] -> Predicted Class: virginica

```






Figure 4: k-NN Confusion Matrix

## Result

Successfully implemented the k-Nearest Neighbour (k-NN) algorithm on the Iris dataset. The correct and wrong predictions were printed, and the classification report showed the precision, recall, and F1-scores for all classes.

# Experiment 7

## Locally Weighted Regression (LWR)

### Aim

Implement the Locally Weighted Regression (LWR) algorithm to fit data points using the ‘tips’ dataset. Visualize the fitted curve and analyze the results.

### Algorithm

1. **Import Required Libraries:** Import the necessary libraries
2. **Load the Dataset:** Load the ‘tips’ dataset, and extract the independent variable (‘total\_bill’) and dependent variable (‘tip’).
3. **Define Kernel Function:** Define a kernel function (e.g., Gaussian kernel) that will compute weights for data points based on their distance from the query point.
4. **Weighted Linear Regression:** For each query point:
  - Compute the weights using the kernel function.
  - Perform weighted linear regression to calculate the predicted value for the query point.
5. **Make Predictions:** Repeat the prediction for a range of query points, which will create a smooth fitted curve.
6. **Visualize Results:** Visualize the original data points along with the fitted curve on a scatter plot.

# Implementation

```

1 # Importing necessary libraries
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6
7 # Load the tips dataset
8 tips = sns.load_dataset('tips')
9 X = tips['total_bill'].values
10 y = tips['tip'].values
11 X = X.reshape(-1, 1)
12
13 # Define Gaussian kernel
14 def kernel(x, x_point, tau):
15     return np.exp(-np.sum((x - x_point) ** 2, axis=1) / (2 * tau ** 2))
16
17 # Locally Weighted Regression function
18 def locally_weighted_regression(x_query, X, y, tau):
19     m = len(X)
20     weights = kernel(X, x_query, tau)
21     W = np.diag(weights)
22     X_aug = np.hstack((np.ones((m, 1)), X))
23     theta = np.linalg.pinv(X_aug.T @ W @ X_aug) @ (X_aug.T @ W @ y)
24     return np.dot([1, x_query], theta)
25
26 # Perform predictions for a range of x values
27 tau = 5
28 x_values = np.linspace(X.min(), X.max(), 200)
29 y_pred = np.array([locally_weighted_regression(np.array([x]), X, y, tau) \
30     for x in x_values])
31
32 # Plot original data and fitted curve
33 plt.figure(figsize=(10, 6))
34 plt.scatter(X, y, color='blue', alpha=0.5, label='Original Data')
35 plt.plot(x_values, y_pred, color='red', label='LWR Fit (tau=5)', linewidth=2)
36 plt.xlabel('Total Bill')
37 plt.ylabel('Tip')
38 plt.legend()
39 plt.show()

```

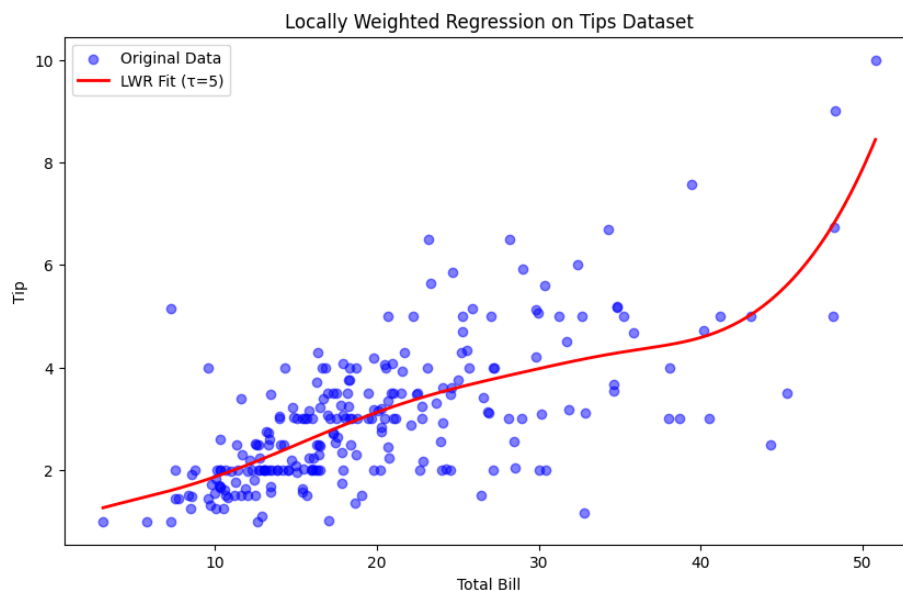


Figure 5: Fitted Curve for Locally Weighted Regression (LWR) on Tips Dataset

## Result

Successfully implemented the Locally Weighted Regression (LWR) algorithm on the 'tips' dataset.

# Experiment 8

## 5-Fold Cross Validation

### Aim

The aim of this experiment is to perform 5-fold cross-validation on the Breast Cancer dataset using Logistic Regression. Evaluate the model's performance based on the metrics of accuracy, precision, recall, and F1-score.

### Algorithm

1. **Import Required Libraries:** Import the necessary libraries
2. **Load the Dataset:** Load the Breast Cancer dataset from the 'sklearn.datasets' module.
3. **Initialize the Classifier:** Use Logistic Regression as the classifier.
4. **Set Up Cross-Validation:** Implement 5-fold cross-validation using the 'StratifiedKFold' method to preserve the percentage of samples for each class in each fold.
5. **Model Evaluation:** Use cross-validation to evaluate the classifier's performance based on accuracy, precision, recall, and F1-score for each fold.
6. **Plot the Results:** Plot the evaluation metrics (accuracy, precision, recall, F1-score) for each fold to visualize the performance.

# Implementation

```

1 from sklearn.datasets import load_breast_cancer
2 from sklearn.model_selection import cross_validate, StratifiedKFold
3 from sklearn.linear_model import LogisticRegression
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7
8 # Load the Breast Cancer dataset
9 data = load_breast_cancer()
10 X = data.data
11 y = data.target
12
13 # Initialize the classifier (you can change it to any classifier)
14 model = LogisticRegression(max_iter=10000, random_state=42)
15
16 # Set up 5-fold cross-validation
17 kf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
18
19 # Cross-validate the model using 5-fold cross-validation
20 cv_results = cross_validate(
21     model, X, y, cv=kf, scoring=['accuracy', 'precision', 'recall', 'f1'], \
22     return_train_score=False
23 )
24
25 # Extract metrics for each fold
26 accuracy = cv_results['test_accuracy']
27 precision = cv_results['test_precision']
28 recall = cv_results['test_recall']
29 f1 = cv_results['test_f1']
30
31 # Display the results for each fold
32 results_df = pd.DataFrame({
33     'Fold': np.arange(1, 6),
34     'Accuracy': accuracy,
35     'Precision': precision,
36     'Recall': recall,
37     'F-Score': f1
38 })
39
40 print("Cross-validation results for each fold:")
41 print(results_df)
42
43 # Display average performance metrics across all folds
44 print("\nAverage Performance across all 5 folds:")
45 print(f"Average Accuracy: {accuracy.mean():.4f}")
46 print(f"Average Precision: {precision.mean():.4f}")
47 print(f"Average Recall: {recall.mean():.4f}")
48 print(f"Average F-Score: {f1.mean():.4f}")
49
50 # Plotting the metrics for each fold
51 fig, ax = plt.subplots(figsize=(10, 6))
52
53 # Define the positions of the bars for each fold
54 x = np.arange(5)
55
56 # Bar width
57 width = 0.2

```

```

58
59 # Plotting each metric
60 ax.bar(x - width*1.5, accuracy, width, label='Accuracy', color='blue')
61 ax.bar(x - width*0.5, precision, width, label='Precision', color='green')
62 ax.bar(x + width*0.5, recall, width, label='Recall', color='orange')
63 ax.bar(x + width*1.5, f1, width, label='F1 Score', color='red')
64
65 # Adding labels and title
66 ax.set_xlabel('Fold')
67 ax.set_ylabel('Score')
68 ax.set_title('Evaluation Metrics for Each Fold (5-Fold Cross Validation)')
69 ax.set_xticks(x)
70 ax.set_xticklabels([f'Fold {i}' for i in range(1, 6)])
71 ax.legend()
72
73 # Show the plot
74 plt.tight_layout()
75 plt.show()

```




Figure 6: Evaluation Metrics for Each Fold (5-Fold Cross Validation)

## Result

The 5-fold cross-validation results for the Breast Cancer dataset are shown below, detailing the accuracy, precision, recall, and F1-score for each fold. The average performance metrics across all folds are also presented.

# Experiment 9

## Support Vector Machine for CIFAR-10 Classification

### Aim

The aim of this experiment is to implement Support Vector Machine (SVM) classifier on the CIFAR-10 dataset, using KNN and a 3-layer neural network.

### Algorithm

1. **Import Libraries:** Import necessary libraries (e.g., `numpy`, `sklearn`, `tensorflow`) for data handling and model building.
2. **Load CIFAR-10 Dataset:** Load the CIFAR-10 dataset, normalize pixel values to `[0, 1]`, and flatten images for KNN.
3. **Train KNN Classifier:** Train a KNN classifier with a subset of the dataset.
4. **Predict Using KNN:** Predict labels for a subset of the test data using the trained KNN classifier.
5. **Evaluate KNN:** Evaluate the KNN classifier by calculating accuracy and generating a classification report.
6. **Define Neural Network:** Define a 3-layer neural network with ReLU activations and softmax output for 10 classes.
7. **Train Neural Network:** Train the neural network on a subset of the CIFAR-10 dataset with validation monitoring.
8. **Predict Using Neural Network:** Predict labels for a subset of the test data using the trained neural network.



**9. Evaluate Neural Network:** Evaluate the neural network on the test set and generate a classification report.

**10. Compare Results:** Compare the accuracy and performance of KNN and the neural network.

## Implementation

```

1 import numpy as np
2 from tensorflow.keras.datasets import cifar10
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import accuracy_score, classification_report
5 from sklearn.preprocessing import StandardScaler
6 import tensorflow as tf
7 from sklearn.neighbors import KNeighborsClassifier
8
9 # Load CIFAR-10 dataset
10 (X_train, y_train), (X_test, y_test) = cifar10.load_data()
11
12 # Normalize pixel values to [0, 1]
13 X_train = X_train.astype('float32') / 255.0
14 X_test = X_test.astype('float32') / 255.0
15
16 # Flatten images for KNN and scaling
17 X_train_flat = X_train.reshape(X_train.shape[0], -1)
18 X_test_flat = X_test.reshape(X_test.shape[0], -1)
19
20 # Reshape labels to 1D
21 y_train = y_train.ravel()
22 y_test = y_test.ravel()
23
24 print("Data loaded and preprocessed.")
25
26 # Initialize KNN with k=15 and Manhattan distance
27 knn = KNeighborsClassifier(n_neighbors=15, metric='manhattan')
28
29 # Train KNN on the CIFAR-10 training data
30 print("Training KNN...")
31 knn.fit(X_train_flat[:20000], y_train[:20000]) # Use a subset for computational efficiency
32 print("KNN training complete.")
33
34 # Predict on the test set
35 y_pred_knn = knn.predict(X_test_flat[:5000]) # Use a subset for faster testing
36
37 # Evaluate KNN
38 knn_accuracy = accuracy_score(y_test[:5000], y_pred_knn)
39 print(f"KNN Accuracy: {knn_accuracy:.4f}")
40 print("\nClassification Report:")
41 print(classification_report(y_test[:5000], y_pred_knn))
42
43 # Define the 3-layer neural network
44 model = tf.keras.Sequential([
45     tf.keras.layers.Flatten(input_shape=(32, 32, 3)), # Input layer
46     tf.keras.layers.Dense(256, activation='relu'), # First hidden layer
47     tf.keras.layers.Dense(128, activation='relu'), # Second hidden layer
48     tf.keras.layers.Dense(10, activation='softmax') # Output layer for 10 classes

```

```
49 ])
```

```
50
```

```
51 # Compile the model
```

```
52 model.compile(optimizer='adam',
```

```
53               loss='sparse_categorical_crossentropy',
```

```
54               metrics=['accuracy'])
```

```
55
```

```
56 # Train the model
```

```
57 print("Training 3-layer neural network...")
```

```
58 history = model.fit(X_train[:20000], y_train[:20000], epochs=10,
```

```
59                   validation_data=(X_test[:5000], y_test[:5000]),
```

```
60                   batch_size=64)
```

```
61 print("Neural network training complete.")
```

```
62
```

```
63 # Evaluate the model on the test set
```

```
64 test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=2)
```

```
65 print(f"Neural Network Test Accuracy: {test_accuracy:.4f}")
```

```
66
```

```
67 # Predict on the test set
```

```
68 y_pred_nn = np.argmax(model.predict(X_test), axis=1)
```

```
69
```

```
70 # Classification report for Neural Network
```

```
71 print("\nClassification Report for Neural Network:")
```

```
72 print(classification_report(y_test, y_pred_nn))
```




Figure 7: Prediction result

## Result

Successfully implemented SVM using two machine learning models, a K-Nearest Neighbors (KNN) classifier and a 3-layer neural network, to classify images from the CIFAR-10 dataset.

# Experiment 10

## Artificial Neural Network with Backpropagation

### Aim

To build an Artificial Neural Network (ANN) using the Backpropagation algorithm and test it with appropriate datasets.

### Algorithm

1. Import necessary libraries
2. Initialize the network parameters, including weights, biases, and hyperparameters (learning rate, number of epochs, etc.).
3. Define the activation function and its derivative (e.g., sigmoid).
4. Perform forward propagation:
  - Compute the output of each layer using the activation function.
  - Calculate the final output.
5. Compute the error using a loss function (e.g., Mean Squared Error).
6. Perform backpropagation:
  - Calculate the gradient of the loss with respect to each parameter (weights and biases).
  - Update weights and biases using gradient descent.
7. Repeat steps 3 to 5 for the specified number of epochs.
8. Test the network with unseen data and analyze its performance.

# Implementation

```

1 # Python code to implement the Artificial Neural Network with Backpropagation
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 # Sigmoid activation function and its derivative
6 def sigmoid(x):
7     return 1 / (1 + np.exp(-x))
8
9 def sigmoid_derivative(x):
10    return x * (1 - x)
11
12 # Dataset (XOR Problem)
13 X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
14 y = np.array([[0], [1], [1], [0]])
15
16 # Initialize hyperparameters
17 input_neurons = 2      # Number of input features
18 hidden_neurons = 4     # Number of hidden layer neurons
19 output_neurons = 1     # Number of output neurons
20 learning_rate = 0.5    # Learning rate
21 epochs = 10000         # Number of iterations
22
23 # Initialize weights and biases
24 np.random.seed(42)    # For reproducibility
25 weights_input_hidden = np.random.uniform(size=(input_neurons, hidden_neurons))
26 weights_hidden_output = np.random.uniform(size=(hidden_neurons, output_neurons))
27 bias_hidden = np.random.uniform(size=(1, hidden_neurons))
28 bias_output = np.random.uniform(size=(1, output_neurons))
29
30 # Store loss for plotting
31 loss_history = []
32
33 # Training the ANN
34 for epoch in range(epochs):
35     # Forward propagation
36     hidden_input = np.dot(X, weights_input_hidden) + bias_hidden
37     hidden_output = sigmoid(hidden_input)
38     final_input = np.dot(hidden_output, weights_hidden_output) + bias_output
39     final_output = sigmoid(final_input)
40
41 # Compute error
42 error = y - final_output
43 loss = np.mean(np.square(error)) # Mean squared error
44 loss_history.append(loss)
45
46 # Backpropagation
47 d_output = error * sigmoid_derivative(final_output)
48 error_hidden_layer = d_output.dot(weights_hidden_output.T)
49 d_hidden_layer = error_hidden_layer * sigmoid_derivative(hidden_output)
50
51 # Update weights and biases
52 weights_hidden_output += hidden_output.T.dot(d_output) * learning_rate
53 weights_input_hidden += X.T.dot(d_hidden_layer) * learning_rate
54 bias_hidden += np.sum(d_hidden_layer, axis=0, keepdims=True) * learning_rate
55 bias_output += np.sum(d_output, axis=0, keepdims=True) * learning_rate
56
57 # Print loss every 1000 epochs

```

```

58     if epoch % 1000 == 0:
59         print(f"Epoch {epoch}, Loss: {loss:.4f}")
60
61 # Testing the ANN
62 print("\nTesting the ANN:")
63 for x, target in zip(X, y):
64     hidden_input = np.dot(x, weights_input_hidden) + bias_hidden
65     hidden_output = sigmoid(hidden_input)
66     final_input = np.dot(hidden_output, weights_hidden_output) + bias_output
67     final_output = sigmoid(final_input)
68     print(f"Input: {x}, Predicted: {final_output.round(2)}, Target: {target}")
69
70 # Plotting the loss curve
71 plt.figure(figsize=(10, 5))
72 plt.plot(loss_history, label="Training Loss", color="blue")
73 plt.title("Loss Curve")
74 plt.xlabel("Epochs")
75 plt.ylabel("Mean Squared Error")
76 plt.legend()
77 plt.grid()
78 plt.show()
79
80 # Plotting the decision boundary
81 def plot_decision_boundary(X, y, model_predict):
82     x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
83     y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
84     xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),
85                          np.arange(y_min, y_max, 0.1))
86     grid = np.c_[xx.ravel(), yy.ravel()]
87     predictions = model_predict(grid)
88     predictions = predictions.reshape(xx.shape)
89     plt.contourf(xx, yy, predictions, alpha=0.8, cmap=plt.cm.Spectral)
90     plt.scatter(X[:, 0], X[:, 1], c=y.flatten(), edgecolor='k', cmap=plt.cm.Spectral)
91     plt.title("Decision Boundary")
92     plt.xlabel("Feature 1")
93     plt.ylabel("Feature 2")
94     plt.show()
95
96 # Define model prediction function
97 def model_predict(data):
98     hidden_input = np.dot(data, weights_input_hidden) + bias_hidden
99     hidden_output = sigmoid(hidden_input)
100    final_input = np.dot(hidden_output, weights_hidden_output) + bias_output
101    return (sigmoid(final_input) > 0.5).astype(int)
102
103 # Plot the decision boundary
104 plot_decision_boundary(X, y, model_predict)
105
106 #output
107 Epoch 0, Loss: 0.3855
108 Epoch 1000, Loss: 0.0168
109 Epoch 2000, Loss: 0.0025
110 Epoch 3000, Loss: 0.0012
111 Epoch 4000, Loss: 0.0008
112 Epoch 5000, Loss: 0.0006
113 Epoch 6000, Loss: 0.0004
114 Epoch 7000, Loss: 0.0004
115 Epoch 8000, Loss: 0.0003
116 Epoch 9000, Loss: 0.0003
117
118 Testing the ANN:

```

```

119 Input: [0 0], Predicted: [[0.02]], Target: [0]
120 Input: [0 1], Predicted: [[0.98]], Target: [1]
121 Input: [1 0], Predicted: [[0.99]], Target: [1]
122 Input: [1 1], Predicted: [[0.01]], Target: [0]

```




Figure 8: Loss Curve




Figure 9: Loss Curve

## **Result**

The Artificial Neural Network was successfully implemented using the Backpropagation algorithm. The network was trained on the XOR dataset and achieved accurate predictions.